

Description

For assignment 4, we're looking for you to be able to write software that *works with* a file system.

Specifically, you're going to be writing code that can read an exFAT file system, and extract files from that file system.

Assignment 4 is due Friday, April 22nd, 2021 at 11:59pm (Winnipeg time).

General submission requirements

Submissions that do not follow these requirements will be rejected (i.e., you will receive a score of 0).

- All solutions must be compiled and run on a Linux system.
- All solutions must be compiled by issuing the command `make` in the directory containing the `Makefile`. No other build commands are acceptable, even if documented in the `README.md`.
 - Forget how to make a `README.md`? Never knew how to make `README.md`? Thankfully, you can go to <https://readme.so/> and build a nice, Markdown-formatted `README.md` with a nice preview and some handy buttons to help you build it.
- All solutions must *run* to successful completion. Premature termination for



any reason (no obvious output, Segmentation Fault, Bus Error, etc.) is considered to be a solution implementation that does not run.

Code that runs for an unreasonably long time (e.g., > 30 seconds) without terminating is also considered to be a solution implementation that does not run.

- All solutions must *compile*. Code that does not compile will not be evaluated.
- Programs must produce no errors when compiled with **all** of the flags

-Wall -Wpedantic -Wextra -Werror

Note that `-Werror` *prevents* your code from being **compiled** when warnings are present.

If all of these flags are not in your `Makefile`, your submission will be treated as though it

- Your code must be submitted electronically
 - No late submissions are accepted
- humanitoba.ca.
- is enforced

Reminder: All testing tool, along with related similarity problems.

Implement *reader*






Get pulled into reading!

In this question, you will write a program that reads an exFAT-formatted volume. For this assignment you are provided with a couple of volume images, but you are encouraged to read a USB flash, SDHC, or SDXC drive as a raw device (e.g., `/dev/sda3`), provided that you have physical access to a Linux machine.

exFAT documentation

The general structure of the program you should take as a reference is provided in class, so you should take it as a starting point.

You should also read the exFAT specification, which is a comprehensive source of truth for the format.



discussed in class, so you should take it as a starting point.

This document is your universal source of truth for the format.

Volumes

Two volumes have been provided for you. The two volumes are distributed as raw images, describing the two volumes, and the program should read them.

am. The two volumes are distributed as raw images, describing the two volumes, and the program should read them.

The graders may not use these volumes to evaluate your work.

Required commands

You should implement 3 commands that are all part of the same program: `info`, `list`, and `get`.

info

The `info` command will print information about the volume. The command, assuming your program is named `exfat`, would be

```
./exfat imagename info
```

Print out the following:

- Volume label.
- Volume serial number.
- Free space on the volume in KB.
- The cluster size, both in sectors and in bytes OR KB.

Note: 1KB ↔ 1024 bytes.

Volume label

The volume label is encoded as a directory entry in the root directory. You can find information about the volume label in section [7.3 Volume Label Directory Entry](#).

The Volume label is a Unicode-formatted string, please see [Unicode and ASCII](#) at the bottom for some additional information.

Volume serial

The volume serial number is encoded as a directory entry in the root directory. You can find information about it both in section [3.1.11 Volume Serial Number](#) and in section [7.4 Volume Serial Number Directory Entry](#).

You should print the volume serial number.

Free space

exFAT uses allocation units (clusters) to store data. You can find information about the free space in the volume (the number of free clusters) in section [3.1.12 Free Space](#) and in section [7.5 Free Space Directory Entry](#).

The allocation bitmap is encoded as a directory entry in the root directory. You can find information about the allocation bitmap in section [7.1 Allocation Bitmap Directory Entry](#).

You should use code that you wrote for lab 4 to determine this value.

Cluster size

The cluster size in sectors and the sector size in bytes can be found in the boot sector. You can find more information about both of these fields in [3.1 Main and Backup Boot Sector Sub-regions](#).

Remember: These fields are called __Shift because you can use the left shift operator (<<) to quickly compute powers of two:

```
uint8_t x = 0x1 << 3; // 2^3 -> 8
```

list

The `list` command will recursively print all files and directories in the volume. The output should look *roughly* like the output from `tree` (try running `tree $YOUR_A2` on `aviary.cs.umanitoba.ca`).

You don't have to print out fancy Extended ASCII characters, but instead can use the `-` symbol to denote depth, where the number of `-` characters indicates the depth of the file/folder. Using the Extended ASCII set would be a nice touch, though.

The command, assuming your program is named `exfat`, would be

```
./exfat imagename list
```

Output should look something like this:

```
egret.cs.umanitoba.ca
```

```
File: file.txt
```

```
Directory: folder1
```

```
- Directory: folder2
```

```
-- Directory: folder3
```

```
--- File: thefile.txt
```

```
File: file2.txt
```

In this example,

file.txt, and

file2.txt, and

to one folder: folder1

folder3. The folder2

Order of files is not important (you can print out more directories, for example). The recommended strategy here is to just print out files and folders in the same order that they are in in terms of the sequence of `DirectoryEntry` in the folders. Likewise, the sequence of directories that you follow is also not important. You're going to have to do this recursively, but it's up to you if you want to do a breadth-first or a depth-first implementation.

get

The `get` command will extract a complete file from the volume. The command, assuming your program is named `exfat`, would be

```
./exfat imagename get path/to/file.txt
```

The path you pass to your program is assumed to be an absolute path (it starts at the root directory), even if it doesn't have a leading `/`.

The path you pass to your program should be written out as a file to the same directory as your program with the same name as the file in the exFAT image. If the example above is executed, you should write the contents of the file at `path/to/file.txt` from the disk image to a file named `file.txt` in the same directory as the `./exfat` binary executable.

Implementation notes

The [exFAT file system specification](#) is fairly comprehensive and generally very good, but some parts of the documentation are not straightforward. Additionally, exFAT supports Unicode characters, and while that's an excellent property of a file system, working with Unicode characters in C isn't a great experience.

Below is some additional information that you can use to help guide you through working with exFAT.

Sectors ↔ clusters

The first chunk of the file system (the boot sector or [Sub-regions](#)) are organized in clusters.

The FATs are also organized in clusters. There can be multiple FATs.

The data heap is organized in clusters. It can be the same size as a sector. It can be different. You should consider the mapping from cluster number to sector number.



Root directory and cluster number offsets

An easily overlooked statement in [5.1 Cluster Heap Sub-region](#) is

Importantly, the first cluster of the Cluster Heap has index two, which directly corresponds to the index of `FatEntry[2]`.

That means that cluster indexing is effectively a 2-based array (what kind of a maniac came up with this?). This means that if, for example, the value in the boot sector for the root directory is 5, then the root directory is actually stored starting at cluster number 3 ($5 - 2 = 3$).

Directory entries

Files, directories, and some FS metadata are encoded as directory entries in the boot sector. The description of the boot sector is in the [exFAT file system specification](#).

data heap region. The description of how to interpret types of directory entries is not great (what the heck does “Benign primary” even mean???)

While you actually can figure out the specific `EntryType` values by building up sequences of bits for each type of `DirectoryEntry`, it’s easier to just look up these values somewhere else. An admittedly spammy looking page selling some software called “Active UNDELETE” has a nice list of the exFAT directory entry types that includes actual numeric values that you can use for comparing against the `EntryType` field in the directory entry.

The directory entries that correspond to one single file will exist as **three** separate, but sequential directory entries:

1. A file directory entry.
2. *Exactly 1* stream extension directory entry.
3. *At least 1* file name directory entry.

You need to read the file name of the file that you’re looking for.

Note: Individual directory entries point to the cluster, but an entry set for a file contains multiple entries, if the cluster chain for the file is not contiguous (e.g. cluster 3 and cluster 10), the file name directory entry points to cluster 3, but the file name directory entry points to cluster 10.

Unicode and

Many fields in the directory entry are Unicode strings. That means that characters are stored in 16-bit units. Provided that your system is little-endian, you can convert a Unicode string to a C string by stripping off the top 8 bits of each character.

For example, if you have a Unicode string containing the Arabic text “أنت ح دث اللغ ة 7”, the C string would be “أنت ح دث اللغ ة 7”. If you have a Unicode string containing the Chinese text “沒有中文”, the C string would be “沒有中文”. If you have a Unicode string containing the Hindi text “बघा पनाघा ठग”, the C string would be “बघा पनाघा ठग”. If you have a Unicode string containing the emoji “👉”, the C string would be “👉”.

```
/**
 * Convert a Unicode-formatted string containing only ASCII characters
 * into a regular ASCII-formatted string (16 bit chars to 8 bit
 * chars).
 *
 * NOTE: this function does a heap allocation for the string it
 * returns, caller is responsible for `free`-ing the allocation
 * when necessary.
 *
 * uint16_t *unicode_string: the Unicode-formatted string to be
 * converted.
 * uint8_t length: the length of the Unicode-formatted string (in
```

```

* characters).
*
* returns: a heap allocated ASCII-formatted string.
*/
static char *unicode2ascii( uint16_t *unicode_string, uint8_t length )
{
    assert( unicode_string != NULL );
    assert( length > 0 );

    char *ascii_string = NULL;

    if ( unicode_string != NULL && length > 0 )
    {
        // +1 for a NULL terminator
        ascii_string = calloc( sizeof(char), length + 1);

        if ( a
        {
            //
            //
            fo
            {
                //
                as
            }
        }
    }

    return ascii_string;
}

```



You are permitted to use this code fragment in your submission.

Boot sector

We provided this to you for lab 4, but here's the complete boot sector struct for your convenience:

```

#pragma pack(1)
#pragma pack(push)
typedef struct MAIN_BOOT_SECTOR
{
    uint8_t jump_boot[3];
    char fs_name[8];

```



```

uint8_t must_be_zero[53];
uint64_t partition_offset;
uint64_t volume_length;
uint32_t fat_offset;
uint32_t fat_length;
uint32_t cluster_heap_offset;
uint32_t cluster_count;
uint32_t first_cluster_of_root_directory;
uint32_t volume_serial_number;
uint16_t fs_revision;
uint16_t fs_flags;
uint8_t bytes_per_sector_shift;
uint8_t sectors_per_cluster_shift;
uint8_t number_of_fats;
uint8_t drive_select;
uint8_t partition;
uint8_t reserved;
uint8_t boot_sector;
uint16_t boot_signature;
} main_boot_sector;

#pragma pack(1)

```

As with lab 4, you will need to use the `main_boot_sector` structure to parse the boot sector.

Binary files

Remember that we are working with binary files. With that in mind, you will need to use the tools you used for parsing an ELF-file. Those tools include

- The `uintX_t` family of integers from `stdint.h`.
- The `hexdump` command-line tool.

In assignment 1 we were working with binary files where the layout of the structures in the file conditionally changed based on the type of file that it was (e.g., 32-bit vs 64-bit files). The exFAT file system **does not** have any structures that are sized conditionally. That means that you can actually use (and are recommended to use) the “read the entire struct” strategy. Remember: this means that you need to use preprocessor directives (`#pragma pack`) to make sure that the compiler doesn’t change the layout of your struct.

Evaluation

Implementation

5 points are awarded for code quality and design:

Level	Description
0	The code is very poor quality (e.g., no comments at all, no functions, poor naming conventions for variables, etc).
1–3	The code is <i>low</i> quality, while some coding standards are applied, their uses is inconsistent (e.g., inconsistent use of comments, <i>some</i> functions but functions might do too much, code is repeated that <i>should</i> be in a function, etc).
4–5	This is the maximum level you can earn if the implementation of your program is substantially incomplete.
Each command	The code is consistently
15 points (5 × 3)	s, for a total of
Level	
0	No att compil
1–2	The su require
3–4	The su are stil
5	The submitted code is complete, all major functionality works as expected.



Submitting your assignment

Submissions will be made using the `handin` command available on all CS UNIX systems.

You should be submitting *at least* the following files:

- A `Makefile`.
- A `README.md` that includes a summary of how to compile and run your program (compiling *should* be “run make”, but you should tell the grader how

to run your program, e.g., what arguments to pass to the program on the command line).

- Your solution for question 1 (probably just 1 .c file).

Please **do not** include any disk images in your submission.

If your files are in a folder named `my_a4`, you would run the command:

```
handin 3430 a4 my_a4
```

If you need more information, you can see `man handin`.


General Advice

Here's some general advice that you can choose to follow or ignore:

- A large green QR code with a white WeChat logo in the center, used for contacting the author on WeChat.









Reflect in ePortfolio



Open with docReader



You have viewed this to

Last Visited Apr 14, 2022 11:3

